

Assignment 2: May the Force...

Group: I Miloncini

Tasks:

- **E) Comparison and benchmark:**

In `data` directory there's a file, `dt_tot.txt`, containing the parameters and the results of the tests.

In the `plots` folder you can find 3 .pdf files:

- `plot_dir_model.pdf` contains the benchmark and comparison between the acceleration estimation functions built with the *direct* method. Data gathered with `pyfalcon` are displayed too. These functions were run for a maximum of 5000 particles, due to the long computational time and expensive computational power. You can see that among the three `acc_dir_*`, `acc_dir_diego` is slightly better than `acc_dir_vepe`. Thus we decided to implement it in the main code as our `acceleration_direct`.
- `plot_vect_model.pdf` contains the benchmark and comparison between the acceleration estimation functions built in a *vectorized* fashion. Data gathered with `pyfalcon` are displayed too. This time we run the simulation till a maximum of 50.000 particles, thanks to the computational power of the **Demoblack** server. You can see that among them, the one with the better performance is `acc_onearray_vepe` (in this case slightly better than `acc_vect_diego`). Thus we decided to implement it in the main code as our `acceleration_direct_vectorized`. The function uses the broadcasting operations of `numpy.array`.
In this function we also implemented the computation of the jerk: the line `jerk_vepe` refers to this case (further considerations below).
- `plot_tot.pdf` contains the comparison between all the models we implemented. It's remarkable the increasing in performance (almost 2 orders) relative to the vectorized models. [`pyfalcon`, using C and memory allocation is playing in another league (it's like Pinerolo vs Roma)]

Speaking about maximum number of particles, `pyfalcon`, based on the [tests](#) run by its developers, for $N \geq 10^5$ the scaling of the CPU time required for the mutual forces becomes essentially linear allowing for a substantial improvement in simulations employing large number of bodies. The only disadvantage is the increased requirement of memory.

Regarding the vectorized functions, a maximum value of N , taking into account a trade-off between computational time and hardware resources (especially RAM, since `numpy.array` requires a lot of it, ~ 150 GB for $N \geq 5 \cdot 10^3$), could be around $5 \cdot 10^4$ (if you are brave enough, and if you have a lot of time and a lot of RAM, you can try also to reach $N = 10^6$).

For non-vectorized functions, the maximum number is even lower. Since the computational time is very high, the maximum number is around $5 \cdot 10^3$ (and necessarily $< 10^4$).

We tried to fit the data we gather (i.e., `dt_tot.txt`) with some functions. The theoretical order of the `acceleration_direct` function should be $O(N^2)$, and we see it from our fitting plot. Regarding the `acceleration_vectorized` function, we expect (looking at the algorithm) that the theoretical order is $O(N^2)$. But understanding how `numpy` deals with data and operations is difficult. In fact we see that the vectorized function is much faster than the direct one, even though both functions are well fitted by a 2nd degree polynomial. Therefore, our conclusion is that at worst the vectorized function goes as $O(N^2)$, but from the results we got we expect that the order is much lower.

- **F) (Optional) Jerk estimate:**

As mentioned above, we implemented the computation of the jerk directly in `acceleration_direct_vectorised`. Since, as you can see from the benchmark, the computation of the jerk is quite computational demanding we implemented a flag `return_jerk` set, as default, `False`. So the performance, for the calculation of the acceleration only, does not decrease.

Setting the flag to `True`, this is the best that the function can do, since we are using the `numpy.array` broadcasting.

Regarding the jerk, we implemented new/updated test functions in `test_acceleration_estimate.py`.

- **Challenge (not mandatory): Need for Speed!**

`acc_onearray_vepe` and `acc_opt_gia` are already the optimized version of the corresponding functions `acc_vect_vepe` and `acc_vect_gia`. The three models of the vectorized function are built in slightly different way, and since the best one, among them, is `acc_onearray_vepe`, probably it's the best result achievable using `numpy.array` and broadcasting.

Thus, we are trying to improve the `acceleration_direct` function using [Cython](#), so as to exploit memory allocation and for loops in a C fashion.